

TensorFlowBasics

December 20, 2017

1 TensorFlow Basics

```
In [75]: import tensorflow as tf
```

1.1 Simple constants

```
In [76]: hello = tf.constant("Hello ")
        world = tf.constant("World")
        with tf.Session() as sess:
            result = sess.run(hello+world)
        print(result)
```

```
b'Hello World'
```

1.2 Operations

```
In [77]: a = tf.constant(10)
        b = tf.constant(20)
        with tf.Session() as sess:
            addition = sess.run(a+b)
            multiplication = sess.run(a*b)
        print(addition)
        print(multiplication)
```

```
30
200
```

1.3 Working with a group of operations

```
In [78]: const = tf.constant(10)
        fill_mat = tf.fill((4,4),10)
        myzeros = tf.zeros((4,4))
        myones = tf.ones((4,4))
        myrandn = tf.random_normal((4,4), mean = 0, stddev = 1.0)
        myrandu = tf.random_uniform((4,4), minval=0, maxval=1)
        my_ops = [const, fill_mat, myzeros, myones, myrandn, myrandu]
```

```

        with tf.Session() as sess:
            for op in my_ops:
                print(op)
                print(sess.run(op))
                print()

Tensor("Const_7:0", shape=(), dtype=int32)
10

Tensor("Fill:0", shape=(4, 4), dtype=int32)
[[10 10 10 10]
 [10 10 10 10]
 [10 10 10 10]
 [10 10 10 10]]

Tensor("zeros:0", shape=(4, 4), dtype=float32)
[[ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]]

Tensor("ones:0", shape=(4, 4), dtype=float32)
[[ 1.  1.  1.  1.]
 [ 1.  1.  1.  1.]
 [ 1.  1.  1.  1.]
 [ 1.  1.  1.  1.]]

Tensor("random_normal:0", shape=(4, 4), dtype=float32)
[[-0.22524352  0.48237228 -0.82078487  3.01135111]
 [ 0.00545869 -0.834593   -0.92740136  0.69436634]
 [-0.65776408 -1.27777672 -0.1419917   0.52560955]
 [ 1.68673515 -1.74032402 -0.78021753 -1.37309158]]

Tensor("random_uniform:0", shape=(4, 4), dtype=float32)
[[ 0.57191324  0.23138583  0.64000297  0.05784225]
 [ 0.70652795  0.35563576  0.8453964   0.2437402 ]
 [ 0.60556591  0.20199621  0.04805243  0.81644404]
 [ 0.12995601  0.93353736  0.22052491  0.71197939]]

```

1.4 Matrix multiplication

```

In [79]: a = tf.constant([[1,2],[3,4]])
         b = tf.constant([[10],[100]])
         result = tf.matmul(a,b)
         with tf.Session() as sess:
             print(sess.run(result))

```

```
[[210]
 [430]]
```

1.5 Graphs

```
In [80]: graph_one = tf.get_default_graph()
graph_two = tf.Graph()
with graph_two.as_default():
    print(graph_two is tf.get_default_graph())
print(graph_one is tf.get_default_graph())
```

```
True
True
```

1.6 Variables and Placeholders

```
In [81]: sess = tf.InteractiveSession()
```

1.6.1 Variables

```
In [82]: my_tensor = tf.random_uniform((4,4),0,1)
my_var = tf.Variable(initial_value= my_tensor)

init = tf.global_variables_initializer()
sess.run(init)

sess.run(my_var)

Out[82]: array([[ 0.84827316,  0.7577678 ,  0.55914938,  0.30990303],
 [ 0.95548916,  0.70689785,  0.23078871,  0.48589516],
 [ 0.71285999,  0.14615691,  0.81334543,  0.25708985],
 [ 0.82702184,  0.70338643,  0.70925617,  0.55189872]], dtype=float32)
```

1.6.2 Placeholders

```
In [83]: ph = tf.placeholder(tf.float32)
```

1.7 More tensorflow Graphs

```
In [84]: import numpy as np
np.random.seed(101)
tf.set_random_seed(101)
"""
Simple calculation example
"""
rand_a = np.random.uniform(0,100,(5,5))
rand_b = np.random.uniform(0,100,(5,1))
```

```

a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
add_op = a + b
mult_op = a * b
with tf.Session() as sess:
    add_result = sess.run(add_op, feed_dict={a:rand_a,b:rand_b})
    mult_result = sess.run(mult_op, feed_dict={a:rand_a,b:rand_b})
print(add_result)
print()
print(mult_result)

```

```

[[ 151.07165527  156.49855042  102.27921295  116.58396149  167.95948792]
 [ 135.45622253   82.76316071  141.42784119  124.22093201   71.06043243]
 [ 113.30171204   93.09214783   76.06819153  136.43911743  154.42727661]
 [  96.7172699    81.83804321  133.83674622  146.38117981  101.10578918]
 [ 122.72680664  105.98292542   59.04463196   67.98310089   72.89292145]]

```

```

[[ 5134.64404297  5674.25          283.12432861  1705.47070312
  6813.83154297]
 [ 4341.8125       1598.26696777  4652.73388672  3756.8293457    988.9463501 ]
 [ 3207.8112793   2038.10290527  1052.77416992  4546.98046875
  5588.11572266]
 [ 1707.37902832   614.02526855  4434.98876953  5356.77734375
  2029.85546875]
 [ 3714.09838867  2806.64379883   262.76763916   747.19854736
  1013.29199219]]

```

1.8 Example Neural Network

```

In [85]: n_features = 10
         n_dense_neurons = 3
         x = tf.placeholder(tf.float32, (None, n_features))
         W = tf.Variable(tf.random_normal([n_features, n_dense_neurons]))
         b = tf.Variable(tf.ones([n_dense_neurons]))
         xW = tf.matmul(x,W)
         z = tf.add(xW, b)
         a = tf.sigmoid(z)
         init = tf.global_variables_initializer()
         with tf.Session() as sess:
             sess.run(init)
             layer_out = sess.run(a, feed_dict={x:np.random.random([1,n_features])})
         print(layer_out)

```

```


[[ 0.9905979  0.99483466  0.40837812]]

```

1.9 Simple Regression Example

```
In [86]: import matplotlib.pyplot as plt
         %matplotlib inline
         """
         Random numbers with some noise
         """
         x_data = np.linspace(0,10,20) + np.random.uniform(-1,2,20)
         y_label = np.linspace(0,10,20) + np.random.uniform(-1,2,20)
         plt.scatter(x_data, y_label)
```

Out[86]: <matplotlib.collections.PathCollection at 0x2aead915f60>



Regression for $y=mx+b$


```
In [87]: m = tf.Variable(np.random.rand(1))
         b = tf.Variable(np.random.rand(1))
         error = 0
         for x,y in zip(x_data, y_label):
             y_hat = m*x + b
             error += (y-y_hat)**2
         #Gradient Descent
         optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.001)
         train = optimizer.minimize(error)
         init = tf.global_variables_initializer()
         with tf.Session() as sess:
             sess.run(init)
```

```

training_steps = 500
for i in range(training_steps):
    sess.run(train)
    final_slope, final_intercept = sess.run([m,b])
x_test = np.linspace(-2,12,20)
y_pred_plot = final_slope*x_test + final_intercept
plt.plot(x_test, y_pred_plot, 'r')
plt.scatter(x_data, y_label)
plt.xlim(-2,12)
plt.ylim(-2,12)

```

Out[87]: (-2, 12)




1.10 Larger Dataset Regression

```

In [88]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
x_data = np.linspace(0.0,10.0,1000000)
noise = np.random.randn(len(x_data))
y_true = (.5 * x_data) + 5 + noise
x_df = pd.DataFrame(data=x_data, columns = ['X Data'])
y_df = pd.DataFrame(data=y_true, columns = ['Y'])
my_data = pd.concat([x_df,y_df],axis=1)
my_data.sample(n=500).plot(kind = 'scatter', x='X Data', y = 'Y')

```

Out[88]: <matplotlib.axes._subplots.AxesSubplot at 0x2b0781fe978>




```
In [89]: batch_size = 8
m = tf.Variable(np.random.randn(1))
b = tf.Variable(np.random.randn(1))
xph = tf.placeholder(tf.float64, [batch_size])
yph = tf.placeholder(tf.float64, [batch_size])
y_model = m * xph + b
error = tf.reduce_sum(tf.square(yph-y_model))
optimizer = tf.train.GradientDescentOptimizer(learning_rate = .001)
train = optimizer.minimize(error)
init = tf.global_variables_initializer()

with tf.Session() as sess:
    sess.run(init)
    batches = 5000
    for i in range(batches):
        rand_ind = np.random.randint(len(x_data), size=batch_size)
        feed = {xph:x_data[rand_ind],yph:y_true[rand_ind]}
        sess.run(train,feed_dict = feed)
    model_m, model_b = sess.run([m,b])

y_hat = x_data * model_m + model_b
my_data.sample(n=500).plot(kind = 'scatter', x='X Data', y = 'Y')
plt.plot(x_data,y_hat,'r')
```

Out[89]: [<matplotlib.lines.Line2D at 0x2b07821eb00>]



1.11 TensorFlow Estimator

```
In [90]: import tensorflow as tf
         feat_cols = [tf.feature_column.numeric_column('x', shape=[1])]
         estimator = tf.estimator.LinearRegressor(feature_columns = feat_cols)

         #Training sets and test sets
         from sklearn.model_selection import train_test_split
         x_train, x_test, y_train, y_test = train_test_split(
             x_data, y_true, test_size = .3, random_state = 101)

         input_func = tf.estimator.inputs.numpy_input_fn({'x':x_train}, y_train,
                                                         batch_size = 8, num_epochs = None,
                                                         shuffle = True)
         train_input_func = tf.estimator.inputs.numpy_input_fn({'x':x_train}, y_train,
                                                                batch_size = 8, num_epochs = 1000,
                                                                shuffle = False)
         test_input_func = tf.estimator.inputs.numpy_input_fn({'x':x_test}, y_test,
                                                              batch_size = 8, num_epochs = 1000,
                                                              shuffle = False)

         estimator.train(input_fn=input_func, steps=1000)

INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Ripti\AppData\Local\Temp\
```



```

INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmppepxu9cm\\model'
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Saving checkpoints for 1 into C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmppepxu9cm\\model
INFO:tensorflow:loss = 587.854, step = 1
INFO:tensorflow:global_step/sec: 436.786
INFO:tensorflow:loss = 17.1739, step = 101 (0.232 sec)
INFO:tensorflow:global_step/sec: 474.863
INFO:tensorflow:loss = 32.2359, step = 201 (0.210 sec)
INFO:tensorflow:global_step/sec: 481.145
INFO:tensorflow:loss = 18.5654, step = 301 (0.211 sec)
INFO:tensorflow:global_step/sec: 465.578
INFO:tensorflow:loss = 14.1275, step = 401 (0.211 sec)
INFO:tensorflow:global_step/sec: 508.848
INFO:tensorflow:loss = 13.2055, step = 501 (0.196 sec)
INFO:tensorflow:global_step/sec: 486.507
INFO:tensorflow:loss = 8.24483, step = 601 (0.208 sec)
INFO:tensorflow:global_step/sec: 457.21
INFO:tensorflow:loss = 13.2243, step = 701 (0.218 sec)
INFO:tensorflow:global_step/sec: 448.581
INFO:tensorflow:loss = 9.94372, step = 801 (0.223 sec)
INFO:tensorflow:global_step/sec: 524.92
INFO:tensorflow:loss = 10.9645, step = 901 (0.191 sec)
INFO:tensorflow:Saving checkpoints for 1000 into C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmppepxu9cm\\model
INFO:tensorflow:Loss for final step: 16.3635.

```

```

Out[90]: <tensorflow.python.estimator.canned.linear.LinearRegressor at 0x2b0643dec18>

```

```

train_metrics = estimator.evaluate(input_fn=train_input_func, steps=1000)
test_metrics = estimator.evaluate(input_fn=test_input_func, steps=1000)

```

```

In [91]: print('TRAINING DATA')
         print(train_metrics)
         print('TEST DATA')
         print(test_metrics)

```

```

TRAINING DATA

```

```

{'average_loss': 1.1141723, 'loss': 8.9133787, 'global_step': 1000}

```

```

TEST DATA

```

```

{'average_loss': 1.1259887, 'loss': 9.0079098, 'global_step': 1000}

```


```

In [92]: new_data = np.linspace(0,10,10)
         input_fn_predict = tf.estimator.inputs.numpy_input_fn({'x':new_data}, shuffle = False)
         predictions = []
         for pred in estimator.predict(input_fn=input_fn_predict):
             predictions.append(pred['predictions'])
         my_data.sample(n=250).plot(kind='scatter',x='X Data', y='Y')
         plt.plot(new_data, predictions, 'r')

```

```
INFO:tensorflow:Restoring parameters from C:\Users\Ripti\AppData\Local\Temp\tmppexpu9cm\model.ck
```

```
Out[92]: [<matplotlib.lines.Line2D at 0x2b07a5e1748>]
```



1.12 Tensorflow Classifications

For when you have categorical data; let's use a binary classification.

```
In [93]: import pandas as pd
diabetes = pd.read_csv('pima-indians-diabetes.csv')
diabetes.head()
```

```
Out[93]:
```

	Number_pregnant	Glucose_concentration	Blood_pressure	Triceps	Insulin	\
0	6	0.743719	0.590164	0.353535	0.000000	
1	1	0.427136	0.540984	0.292929	0.000000	
2	8	0.919598	0.524590	0.000000	0.000000	
3	1	0.447236	0.540984	0.232323	0.111111	
4	0	0.688442	0.327869	0.353535	0.198582	

	BMI	Pedigree	Age	Class	Group
0	0.500745	0.234415	50	1	B
1	0.396423	0.116567	31	0	C
2	0.347243	0.253629	32	1	B
3	0.418778	0.038002	21	0	B
4	0.642325	0.943638	33	1	C

```

In [94]: # Normalization
cols_to_norm = ['Number_pregnant', 'Glucose_concentration',
                'Blood_pressure', 'Triceps',
                'Insulin', 'BMI', 'Pedigree']
diabetes[cols_to_norm] = diabetes[cols_to_norm].apply(
    lambda x: (x-x.min())/(x.max()-x.min()))

In [95]: # Continuous features
num_preg = tf.feature_column.numeric_column('Number_pregnant')
plasma_gluc = tf.feature_column.numeric_column('Glucose_concentration')
dias_press = tf.feature_column.numeric_column('Blood_pressure')
tricep = tf.feature_column.numeric_column('Triceps')
insulin = tf.feature_column.numeric_column('Insulin')
bmi = tf.feature_column.numeric_column('BMI')
diabetes_pedigree = tf.feature_column.numeric_column('Pedigree')
age = tf.feature_column.numeric_column('Age')

# Categorical features
assigned_group = tf.feature_column.categorical_column_with_vocabulary_list(
    'Group', ['A', 'B', 'C', 'D'])


```

1.12.1 Converting Continuous to Categorical

```

In [96]: import matplotlib.pyplot as plt
%matplotlib inline
diabetes['Age'].hist(bins=20)
age_buckets = tf.feature_column.bucketized_column(
    age, boundaries=[20,30,40,50,60,70,80])
feat_cols = [num_preg ,plasma_gluc,dias_press ,tricep ,
             insulin,bmi,diabetes_pedigree ,assigned_group, age_buckets]

```



1.12.2 Training and Test Sets

```
In [97]: from sklearn.model_selection import train_test_split
x_data = diabetes.drop('Class',axis=1)
labels = diabetes['Class']
X_train, X_test, y_train, y_test = train_test_split(x_data,labels,
                                                    test_size=0.33, random_state=101)
input_func = tf.estimator.inputs.pandas_input_fn(x=X_train,y=y_train,
                                                  batch_size=10,num_epochs=1000,shuffle=True)
model = tf.estimator.LinearClassifier(feature_columns=feat_cols,n_classes=2)
model.train(input_fn=input_func,steps=1000)
```

INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Ripti\AppData\Local\Temp\...
INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmp05nr32m...
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Saving checkpoints for 1 into C:\Users\Ripti\AppData\Local\Temp\tmp05nr32m5\model...
INFO:tensorflow:loss = 6.93147, step = 1
INFO:tensorflow:global_step/sec: 127.351
INFO:tensorflow:loss = 5.57055, step = 101 (0.785 sec)
INFO:tensorflow:global_step/sec: 149.08
INFO:tensorflow:loss = 5.15163, step = 201 (0.671 sec)
INFO:tensorflow:global_step/sec: 155.834
INFO:tensorflow:loss = 4.82526, step = 301 (0.643 sec)
INFO:tensorflow:global_step/sec: 148.857
INFO:tensorflow:loss = 6.45611, step = 401 (0.671 sec)

```

INFO:tensorflow:global_step/sec: 145.597
INFO:tensorflow:loss = 5.20896, step = 501 (0.687 sec)
INFO:tensorflow:global_step/sec: 143.09
INFO:tensorflow:loss = 6.39312, step = 601 (0.700 sec)
INFO:tensorflow:global_step/sec: 148.856
INFO:tensorflow:loss = 5.7191, step = 701 (0.671 sec)
INFO:tensorflow:global_step/sec: 142.274
INFO:tensorflow:loss = 4.20419, step = 801 (0.703 sec)
INFO:tensorflow:global_step/sec: 144.124
INFO:tensorflow:loss = 5.41011, step = 901 (0.696 sec)
INFO:tensorflow:Saving checkpoints for 1000 into C:\Users\Ripti\AppData\Local\Temp\tmp05nr32m5\m
INFO:tensorflow:Loss for final step: 3.82677.

```

```

Out[97]: <tensorflow.python.estimator.canned.linear.LinearClassifier at 0x2b06d8ca3c8>

```

1.12.3 Evaluating the Model

```

In [98]: eval_input_func = tf.estimator.inputs.pandas_input_fn(
        x=X_test,
        y=y_test,
        batch_size=10,
        num_epochs=1,
        shuffle=False)
        results = model.evaluate(eval_input_func)
        results

WARNING:tensorflow:Casting <dtype: 'float32'> labels to bool.
WARNING:tensorflow:Casting <dtype: 'float32'> labels to bool.
INFO:tensorflow:Starting evaluation at 2017-12-21-02:40:17
INFO:tensorflow:Restoring parameters from C:\Users\Ripti\AppData\Local\Temp\tmp05nr32m5\model.ck
INFO:tensorflow:Finished evaluation at 2017-12-21-02:40:18
INFO:tensorflow:Saving dict for global step 1000: accuracy = 0.732283, accuracy_baseline = 0.657

```

```

Out[98]: {'accuracy': 0.73228347,
          'accuracy_baseline': 0.65748036,
          'auc': 0.78271043,
          'auc_precision_recall': 0.61937535,
          'average_loss': 0.53361279,
          'global_step': 1000,
          'label/mean': 0.34251967,
          'loss': 5.2129865,
          'prediction/mean': 0.35916343}

```

```

In [99]: # Predictions
        pred_input_func = tf.estimator.inputs.pandas_input_fn(
            x=X_test,
            batch_size=10,

```

```

        num_epochs=1,
        shuffle=False)
predictions = model.predict(pred_input_func)

```

Let's use a Dense Neural Network

```

In [100]: embedded_group_col = tf.feature_column.embedding_column(
        assigned_group, dimension=4)
feat_cols = [num_preg ,plasma_gluc,dias_press ,tricep ,insulin,
             bmi,diabetes_pedigree ,embedded_group_col, age_buckets]
input_func = tf.estimator.inputs.pandas_input_fn(
        x=X_train,y=y_train,batch_size=10,num_epochs=1000,shuffle=True)

dnn_model = tf.estimator.DNNClassifier(hidden_units=[10,10,10],
        feature_columns=feat_cols,n_classes=2)
dnn_model.train(input_fn=input_func,steps=1000)

```

```

INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Ripti\AppData\Local\Temp\
INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmpbv82dg_
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Saving checkpoints for 1 into C:\Users\Ripti\AppData\Local\Temp\tmpbv82dg_3\mode
INFO:tensorflow:loss = 7.01344, step = 1
INFO:tensorflow:global_step/sec: 131.478
INFO:tensorflow:loss = 5.80182, step = 101 (0.761 sec)
INFO:tensorflow:global_step/sec: 142.598
INFO:tensorflow:loss = 6.48086, step = 201 (0.700 sec)
INFO:tensorflow:global_step/sec: 147.34
INFO:tensorflow:loss = 7.30843, step = 301 (0.679 sec)
INFO:tensorflow:global_step/sec: 142.884
INFO:tensorflow:loss = 6.47545, step = 401 (0.700 sec)
INFO:tensorflow:global_step/sec: 138.675
INFO:tensorflow:loss = 4.0591, step = 501 (0.723 sec)
INFO:tensorflow:global_step/sec: 144.264
INFO:tensorflow:loss = 2.52048, step = 601 (0.693 sec)
INFO:tensorflow:global_step/sec: 152.498
INFO:tensorflow:loss = 3.87574, step = 701 (0.654 sec)
INFO:tensorflow:global_step/sec: 142.214
INFO:tensorflow:loss = 3.00031, step = 801 (0.705 sec)
INFO:tensorflow:global_step/sec: 147.312
INFO:tensorflow:loss = 2.53928, step = 901 (0.678 sec)
INFO:tensorflow:Saving checkpoints for 1000 into C:\Users\Ripti\AppData\Local\Temp\tmpbv82dg_3\m
INFO:tensorflow:Loss for final step: 4.76311.

```

```

Out[100]: <tensorflow.python.estimator.canned.dnn.DNNClassifier at 0x2b07a4cc6a0>

```

```

In [101]: eval_input_func = tf.estimator.inputs.pandas_input_fn(
        x = X_test, y = y_test, batch_size = 10, num_epochs = 1, shuffle = False)

```

```
In [102]: dnn_model.evaluate(eval_input_func)
```

```
WARNING:tensorflow: Casting <dtype: 'float32'> labels to bool.
```

```
WARNING:tensorflow: Casting <dtype: 'float32'> labels to bool.
```

```
INFO:tensorflow: Starting evaluation at 2017-12-21-02:40:30
```

```
INFO:tensorflow: Restoring parameters from C:\Users\Ripti\AppData\Local\Temp\tmpbv82dg_3\model.chk
```

```
INFO:tensorflow: Finished evaluation at 2017-12-21-02:40:31
```

```
INFO:tensorflow: Saving dict for global step 1000: accuracy = 0.751969, accuracy_baseline = 0.657
```

```
Out[102]: {'accuracy': 0.7519685,  
          'accuracy_baseline': 0.65748036,  
          'auc': 0.81444001,  
          'auc_precision_recall': 0.68500859,  
          'average_loss': 0.50868565,  
          'global_step': 1000,  
          'label/mean': 0.34251967,  
          'loss': 4.9694676,  
          'prediction/mean': 0.37802684}
```

2 Saving and Loading Models

```
In [103]: import numpy as np  
import tensorflow as tf  
import matplotlib.pyplot as plt  
%matplotlib inline  
  
np.random.seed(101)  
tf.set_random_seed(101)  
  
# Artificial Data  
x_data = np.linspace(0,10,10) + np.random.uniform(-1.5,1.5,10)  
y_label = np.linspace(0,10,10) + np.random.uniform(-1.5,1.5,10)  
m = tf.Variable(np.random.rand(1))  
b = tf.Variable(np.random.rand(1))  
error = tf.reduce_mean(y_label - (m*x_data+b))  
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.001)  
train = optimizer.minimize(error)  
init = tf.global_variables_initializer()
```

2.1 Saving the Model

```
In [104]: saver = tf.train.Saver()  
with tf.Session() as sess:  
    sess.run(init)  
    epochs = 100  
    for i in range(epochs):  
        sess.run(train)
```

```

# Fetch Back Results
final_slope , final_intercept = sess.run([m,b])
saver.save(sess,'new_models/model_from_course.ckpt')


x_test = np.linspace(-1,11,10)
y_pred_plot = final_slope*x_test + final_intercept

plt.plot(x_test,y_pred_plot,'r')

plt.plot(x_data,y_label,'*')

```

Out[104]: [<matplotlib.lines.Line2D at 0x2aeaec77f0>]



2.2 Loading a Model

```

In [105]: with tf.Session() as sess:
# Restore the model
saver.restore(sess,'new_models/model_from_course.ckpt')
# Fetch Back Results
restored_slope , restored_intercept = sess.run([m,b])
x_test = np.linspace(-1,11,10)
y_pred_plot = restored_slope*x_test + restored_intercept

plt.plot(x_test,y_pred_plot,'r')


plt.plot(x_data,y_label,'*')

```



```
INFO:tensorflow:Restoring parameters from new_models/model_from_course.ckpt
```

```
Out[105]: [<matplotlib.lines.Line2D at 0x2b06c820f98>]
```



3 Regression Exercise

3.1 The Data

Import the `cal_housing.csv` file with `pandas`. Separate it into a training (70%) and testing set(30%).

```
In [106]: import pandas as pd
housing = pd.read_csv('cal_housing_clean.csv')
housing.head()
```

```
Out[106]:
```

	housingMedianAge	totalRooms	totalBedrooms	population	households	\
0	41.0	880.0	129.0	322.0	126.0	
1	21.0	7099.0	1106.0	2401.0	1138.0	
2	52.0	1467.0	190.0	496.0	177.0	
3	52.0	1274.0	235.0	558.0	219.0	
4	52.0	1627.0	280.0	565.0	259.0	

	medianIncome	medianHouseValue
0	8.3252	452600.0
1	8.3014	358500.0

2	7.2574	352100.0
3	5.6431	341300.0
4	3.8462	342200.0

```
In [107]: housing.describe().transpose()
```

```
Out[107]:
```

	count	mean	std	min	\
housingMedianAge	20640.0	28.639486	12.585558	1.0000	
totalRooms	20640.0	2635.763081	2181.615252	2.0000	
totalBedrooms	20640.0	537.898014	421.247906	1.0000	
population	20640.0	1425.476744	1132.462122	3.0000	
households	20640.0	499.539680	382.329753	1.0000	
medianIncome	20640.0	3.870671	1.899822	0.4999	
medianHouseValue	20640.0	206855.816909	115395.615874	14999.0000	

	25%	50%	75%	max
housingMedianAge	18.0000	29.0000	37.00000	52.0000
totalRooms	1447.7500	2127.0000	3148.00000	39320.0000
totalBedrooms	295.0000	435.0000	647.00000	6445.0000
population	787.0000	1166.0000	1725.00000	35682.0000
households	280.0000	409.0000	605.00000	6082.0000
medianIncome	2.5634	3.5348	4.74325	15.0001
medianHouseValue	119600.0000	179700.0000	264725.00000	500001.0000

```
In [108]: from sklearn.model_selection import train_test_split
x_data = housing.drop(['medianHouseValue'],axis=1)
y_val = housing['medianHouseValue']
X_train, X_test, y_train, y_test = train_test_split(
    x_data,y_val,test_size=0.3,random_state=101)
```

3.2 Scale the Feature Data

Use sklearn preprocessing to create a MinMaxScaler for the feature data. Fit this scaler only to the training data. Then use it to transform X_test and X_train. Then use the scaled X_test and X_train along with pd.DataFrame to re-create two dataframes of scaled data.

```
In [109]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaler.fit(X_train)
X_train = pd.DataFrame(data=scaler.transform(X_train),
                        columns = X_train.columns,index=X_train.index)
X_test = pd.DataFrame(data=scaler.transform(X_test),
                       columns = X_test.columns,index=X_test.index)
```

3.3 Create Feature Columns

Create the necessary tf.feature_column objects for the estimator. They should all be treated as continuous numeric_columns.

Create the input function for the estimator object. (play around with batch_size and num_epochs)

Create the estimator model. Use a DNNRegressor. Play around with the hidden units!

Train the model for ~1,000 steps. (Later come back to this and train it for more and check for improvement)

```
In [115]: import tensorflow as tf
          age = tf.feature_column.numeric_column('housingMedianAge')
          rooms = tf.feature_column.numeric_column('totalRooms')
          bedrooms = tf.feature_column.numeric_column('totalBedrooms')
          pop = tf.feature_column.numeric_column('population')
          households = tf.feature_column.numeric_column('households')
          income = tf.feature_column.numeric_column('medianIncome')
          feat_cols = [age, rooms, bedrooms, pop, households, income]
          input_func = tf.estimator.inputs.pandas_input_fn(
              x=X_train, y=y_train,
              batch_size=10, num_epochs=1000, shuffle=True)
          model = tf.estimator.DNNRegressor(hidden_units=[6,6,6], feature_columns=feat_cols)
          # model.train(input_fn=input_func, steps=25000)
```

INFO:tensorflow:Using default config.

WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Ripti\AppData\Local\Temp\

INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmpt8ltkpu

Create a prediction input function and then use the .predict method off your estimator model to create a list of predictions on your test data.

Calculate the RMSE

```
In [113]: from sklearn.metrics import mean_squared_error
          predict_input_func = tf.estimator.inputs.pandas_input_fn(
              x=X_test,
              batch_size=10,
              num_epochs=1,
              shuffle=False)
          pred_gen = model.predict(predict_input_func)
          predictions = list(pred_gen)

          final_preds = []
          for pred in predictions:
              final_preds.append(pred['predictions'])
          mean_squared_error(y_test, final_preds)**0.5
```

INFO:tensorflow:Restoring parameters from C:\Users\Ripti\AppData\Local\Temp\tmpjhqgpype\model.ck

Out[113]: 96258.035354416352

4 Classification Exercise

4.1 THE DATA

Read in the census_data.csv data with pandas

```
In [118]: import pandas as pd
          census = pd.read_csv("census_data.csv")
          census.head()
```

Out[118]:

	age	workclass	education	education_num	marital_status	
0	39	State-gov	Bachelors	13	Never-married	
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	
2	38	Private	HS-grad	9	Divorced	
3	53	Private	11th	7	Married-civ-spouse	
4	28	Private	Bachelors	13	Married-civ-spouse	

	occupation	relationship	race	gender	capital_gain	
0	Adm-clerical	Not-in-family	White	Male	2174	
1	Exec-managerial	Husband	White	Male	0	
2	Handlers-cleaners	Not-in-family	White	Male	0	
3	Handlers-cleaners	Husband	Black	Male	0	
4	Prof-specialty	Wife	Black	Female	0	

	capital_loss	hours_per_week	native_country	income_bracket
0	0	40	United-States	<=50K
1	0	13	United-States	<=50K
2	0	40	United-States	<=50K
3	0	40	United-States	<=50K
4	0	40	Cuba	<=50K

TensorFlow won't be able to understand strings as labels, you'll need to use pandas .apply() method to apply a custom function that converts them to 0s and 1s. This might be hard if you aren't very familiar with pandas, so feel free to take a peek at the solutions for this part.

Convert the Label column to 0s and 1s instead of strings.

```
In [119]: census['income_bracket'].unique()
          def label_fix(label):
              if label==' <=50K':
                  return 0
              else:
                  return 1
          census['income_bracket'] = census['income_bracket'].apply(label_fix)
```

4.2 Perform a Train Test Split on the Data

Create the Feature Columns for tf.estimator

Take note of categorical vs continuous values!

Create the tf.feature_columns for the categorical values. Use vocabulary lists or just use hash buckets.

Create the continuous feature_columns for the continuous values using numeric_column

Put all these variables into a single list with the variable name feat_cols

Create Input Function

Batch_size is up to you. But do make sure to shuffle!

Create your model with tf.estimator

Create a LinearClassifier.(If you want to use a DNNClassifier, keep in mind you'll need to create embedded columns out of the categorical feature that use strings, check out the previous lecture on this for more info.)

Train your model on the data, for at least 5000 steps.

```
In [122]: import tensorflow as tf
          from sklearn.model_selection import train_test_split
          x_data = census.drop('income_bracket',axis=1)
          y_labels = census['income_bracket']

          X_train, X_test, y_train, y_test = train_test_split(x_data,y_labels,test_size=0.3,rand

          # Categorical
          gender = tf.feature_column.categorical_column_with_vocabulary_list(
              "gender", ["Female", "Male"])
          occupation = tf.feature_column.categorical_column_with_hash_bucket(
              "occupation", hash_bucket_size=1000)
          marital_status = tf.feature_column.categorical_column_with_hash_bucket(
              "marital_status", hash_bucket_size=1000)
          relationship = tf.feature_column.categorical_column_with_hash_bucket(
              "relationship", hash_bucket_size=1000)
          education = tf.feature_column.categorical_column_with_hash_bucket(
              "education", hash_bucket_size=1000)
          workclass = tf.feature_column.categorical_column_with_hash_bucket(
              "workclass", hash_bucket_size=1000)
          native_country = tf.feature_column.categorical_column_with_hash_bucket(
              "native_country", hash_bucket_size=1000)

          # Continuous
          age = tf.feature_column.numeric_column("age")
          education_num = tf.feature_column.numeric_column("education_num")
          capital_gain = tf.feature_column.numeric_column("capital_gain")
          capital_loss = tf.feature_column.numeric_column("capital_loss")
          hours_per_week = tf.feature_column.numeric_column("hours_per_week")

          feat_cols = [gender,occupation,marital_status,relationship,
                        education,workclass,native_country,age,
                        education_num,capital_gain,capital_loss,hours_per_week]
          input_func=tf.estimator.inputs.pandas_input_fn(
              x=X_train,y=y_train,batch_size=100,num_epochs=None,shuffle=True)
          model = tf.estimator.LinearClassifier(feature_columns=feat_cols)

          # model.train(input_fn=input_func, steps=5000)
```

```
INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Ripti\AppData\Local\Temp\
INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmpfd82o3_
```

4.3 Evaluation

Create a prediction input function. Remember to only support `X_test` data and keep `shuffle=False`.

Use `model.predict()` and pass in your input function. This will produce a generator of predictions, which you can then transform into a list, with `list()`

Create a list of only the `class_ids` key values from the prediction list of dictionaries, these are the predictions you will use to compare against the real `y_test` values.

Import `classification_report` from `sklearn.metrics` and then see if you can figure out how to use it to easily get a full report of your model's performance on the test data.

```
In [121]: from sklearn.metrics import classification_report
          pred_fn = tf.estimator.inputs.pandas_input_fn(
              x=X_test, batch_size=len(X_test), shuffle=False)
          predictions = list(model.predict(input_fn=pred_fn))
          final_preds = []
          for pred in predictions:
              final_preds.append(pred['class_ids'][0])
          print(classification_report(y_test, final_preds))
```

```
INFO:tensorflow:Restoring parameters from C:\Users\Ripti\AppData\Local\Temp\tmpfimvx4k1\model.ckpt
precision    recall  f1-score   support
```

0	0.86	0.92	0.89	7436
1	0.69	0.53	0.60	2333
avg / total	0.82	0.83	0.82	9769