

# Miscellaneous Topics

December 21, 2017

## 1 Deep Nets with TF Abstractions

Let's explore a few of the various abstractions that TensorFlow offers. You can check out the `tf.contrib` documentation for more options.

## 2 The Data

To compare these various abstractions we'll use a dataset easily available from the SciKit Learn library. The data is comprised of the results of a chemical analysis of wines grown in the same region in Italy by three different cultivators. There are thirteen different measurements taken for different constituents found in the three types of wine. We will use the various TF Abstractions to classify the wine to one of the 3 possible labels.

First let's show you how to get the data:

```
In [1]: from sklearn.datasets import load_wine
        wine_data = load_wine()
        print(type(wine_data))
```

```
<class 'sklearn.utils.Bunch'>
```

The data is a `sklearn.utils.Bunch` object, which is very similar to a dictionary.

```
In [2]: wine_data.keys()
```

```
Out[2]: dict_keys(['data', 'target', 'target_names', 'DESCR', 'feature_names'])
```

You can get a full description with `print(wine_data.DESCR)`. For now, let's go ahead and grab the features and the labels for the data.

```
In [3]: feat_data = wine_data['data']
        labels = wine_data['target']
```

### 2.0.1 Train Test Split

As with any machine learning model, you should do some sort of test train split so you can evaluate your model's performance. Because this particular dataset is small, we'll just do a simple 70/30 train test split and we won't have any holdout data set.

Again, we'll use SciKit-Learn here for convenience:

```
In [4]: from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(feat_data,
                                                            labels,
                                                            test_size=0.3,
                                                            random_state=101)
```

### 2.0.2 Scale the Data

With Neural Network models, its important to scale the data, again we can do this easily with SciKit Learn (I promise we'll get to TensorFlow soon!)

```
In [5]: from sklearn.preprocessing import MinMaxScaler
        scaler = MinMaxScaler()
```

Keep in mind we only fit the scaler to the training data, we don't want to assume we'll have knowledge of future test data.

```
In [6]: scaled_x_train = scaler.fit_transform(X_train)
        scaled_x_test = scaler.transform(X_test)
```

## 3 Abstractions

With our data set up, its now time to explore some TensorFlow abstractions! Let's start with the Estimator API, its one the abstractions featured in the official documentation tutorials.

### 3.1 Estimator API

We first start by importing both tensorflow and the estimator API.

```
In [7]: import tensorflow as tf
        from tensorflow import estimator
```

The estimator API can perform both Deep Neural Network Classification and Regression, as well as straight Linear Classification and Linear Regression. You can

```
In [10]: X_train.shape
```

```
Out[10]: (124, 13)
```

```
In [11]: feat_cols = [tf.feature_column.numeric_column("x", shape=[13])]
```

```

In [12]: deep_model = estimator.DNNClassifier(
        hidden_units=[13,13,13],
        feature_columns=feat_cols,
        n_classes=3,
        optimizer=tf.train.GradientDescentOptimizer(
            learning_rate=0.01))

INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Ripti\AppData\Local\Temp\
INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\Ripti\\AppData\\Local\\Temp\\tmplr_iw3h

In [13]: input_fn = estimator.inputs.numpy_input_fn(
        x={'x':scaled_x_train},
        y=y_train,
        shuffle=True,
        batch_size=10,
        num_epochs=5)

In [14]: deep_model.train(input_fn=input_fn,steps=500)

INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Saving checkpoints for 1 into C:\Users\Ripti\AppData\Local\Temp\tmplr_iw3hk\model
INFO:tensorflow:loss = 12.1808, step = 1
INFO:tensorflow:Saving checkpoints for 62 into C:\Users\Ripti\AppData\Local\Temp\tmplr_iw3hk\mod
INFO:tensorflow:Loss for final step: 7.43543.

Out[14]: <tensorflow.python.estimator.canned.dnn.DNNClassifier at 0x1a8c4fd2780>

In [15]: input_fn_eval = estimator.inputs.numpy_input_fn(
        x={'x':scaled_x_test},
        shuffle=False)

In [16]: preds = list(deep_model.predict(input_fn=input_fn_eval))

INFO:tensorflow:Restoring parameters from C:\Users\Ripti\AppData\Local\Temp\tmplr_iw3hk\model.ck

In [17]: predictions = [p['class_ids'][0] for p in preds]

In [18]: from sklearn.metrics import confusion_matrix,classification_report

In [19]: print(classification_report(y_test,predictions))

              precision    recall  f1-score   support

0               1.00        0.53        0.69         19
1               0.62        0.68        0.65         22
2               0.65        1.00        0.79         13

```

avg / total	0.76	0.70	0.70	54
-------------	------	------	------	----

---

---

## 4 TensorFlow Keras

### 4.0.1 Create the Model

```
In [20]: from tensorflow.contrib.keras import models
```

```
In [21]: dnn_keras_model = models.Sequential()
```

### 4.0.2 Add Layers to the model

```
In [22]: from tensorflow.contrib.keras import layers
```

```
In [23]: dnn_keras_model.add(
          layers.Dense(units=13,input_dim=13,activation='relu'))
```

```
In [24]: dnn_keras_model.add(layers.Dense(units=13,activation='relu'))
          dnn_keras_model.add(layers.Dense(units=13,activation='relu'))
```

```
In [25]: dnn_keras_model.add(layers.Dense(units=3,activation='softmax'))
```

### 4.0.3 Compile the Model

```
In [26]: from tensorflow.contrib.keras import losses,optimizers,metrics
```

```
In [27]: # explore these
          # losses.
```

```
In [28]: #optimizers.
```

```
In [29]: losses.sparse_categorical_crossentropy
```

```
Out[29]: <function tensorflow.python.keras._impl.keras.losses.sparse_categorical_crossentropy>
```

```
In [30]: dnn_keras_model.compile(optimizer='adam',
                                loss='sparse_categorical_crossentropy',
                                metrics=['accuracy'])
```

#### 4.0.4 Train Model

```
In [31]: dnn_keras_model.fit(scaled_x_train,y_train,epochs=50)
```

```
Epoch 1/50
124/124 [=====] - 0s - loss: 1.0800 - acc: 0.4274
Epoch 2/50
124/124 [=====] - 0s - loss: 1.0675 - acc: 0.5968
Epoch 3/50
124/124 [=====] - 0s - loss: 1.0561 - acc: 0.5806
Epoch 4/50
124/124 [=====] - 0s - loss: 1.0480 - acc: 0.5806
Epoch 5/50
124/124 [=====] - 0s - loss: 1.0377 - acc: 0.5645
Epoch 6/50
124/124 [=====] - 0s - loss: 1.0279 - acc: 0.5726
Epoch 7/50
124/124 [=====] - 0s - loss: 1.0172 - acc: 0.5726
Epoch 8/50
124/124 [=====] - 0s - loss: 1.0067 - acc: 0.5968
Epoch 9/50
124/124 [=====] - 0s - loss: 0.9947 - acc: 0.6129
Epoch 10/50
124/124 [=====] - 0s - loss: 0.9815 - acc: 0.6452
Epoch 11/50
124/124 [=====] - 0s - loss: 0.9665 - acc: 0.6532
Epoch 12/50
124/124 [=====] - 0s - loss: 0.9494 - acc: 0.6935
Epoch 13/50
124/124 [=====] - 0s - loss: 0.9320 - acc: 0.6935
Epoch 14/50
124/124 [=====] - 0s - loss: 0.9138 - acc: 0.7097
Epoch 15/50
124/124 [=====] - 0s - loss: 0.8937 - acc: 0.7258
Epoch 16/50
124/124 [=====] - 0s - loss: 0.8712 - acc: 0.7984
Epoch 17/50
124/124 [=====] - 0s - loss: 0.8473 - acc: 0.8145
Epoch 18/50
124/124 [=====] - 0s - loss: 0.8225 - acc: 0.8468
Epoch 19/50
124/124 [=====] - 0s - loss: 0.7966 - acc: 0.8710
Epoch 20/50
124/124 [=====] - 0s - loss: 0.7674 - acc: 0.8952
Epoch 21/50
124/124 [=====] - 0s - loss: 0.7387 - acc: 0.8952
Epoch 22/50
124/124 [=====] - 0s - loss: 0.7075 - acc: 0.9032
Epoch 23/50
```

124/124 [=====] - 0s - loss: 0.6763 - acc: 0.9113  
Epoch 24/50  
124/124 [=====] - 0s - loss: 0.6443 - acc: 0.9113  
Epoch 25/50  
124/124 [=====] - 0s - loss: 0.6138 - acc: 0.9194  
Epoch 26/50  
124/124 [=====] - 0s - loss: 0.5814 - acc: 0.9194  
Epoch 27/50  
124/124 [=====] - 0s - loss: 0.5515 - acc: 0.9194  
Epoch 28/50  
124/124 [=====] - 0s - loss: 0.5220 - acc: 0.9274  
Epoch 29/50  
124/124 [=====] - 0s - loss: 0.4942 - acc: 0.9274  
Epoch 30/50  
124/124 [=====] - 0s - loss: 0.4655 - acc: 0.9274  
Epoch 31/50  
124/124 [=====] - 0s - loss: 0.4393 - acc: 0.9274  
Epoch 32/50  
124/124 [=====] - 0s - loss: 0.4140 - acc: 0.9274  
Epoch 33/50  
124/124 [=====] - 0s - loss: 0.3914 - acc: 0.9435  
Epoch 34/50  
124/124 [=====] - 0s - loss: 0.3689 - acc: 0.9435  
Epoch 35/50  
124/124 [=====] - 0s - loss: 0.3485 - acc: 0.9435  
Epoch 36/50  
124/124 [=====] - 0s - loss: 0.3292 - acc: 0.9516  
Epoch 37/50  
124/124 [=====] - 0s - loss: 0.3126 - acc: 0.9516  
Epoch 38/50  
124/124 [=====] - 0s - loss: 0.2964 - acc: 0.9516  
Epoch 39/50  
124/124 [=====] - 0s - loss: 0.2827 - acc: 0.9516  
Epoch 40/50  
124/124 [=====] - 0s - loss: 0.2692 - acc: 0.9516  
Epoch 41/50  
124/124 [=====] - 0s - loss: 0.2569 - acc: 0.9516  
Epoch 42/50  
124/124 [=====] - 0s - loss: 0.2453 - acc: 0.9516  
Epoch 43/50  
124/124 [=====] - 0s - loss: 0.2353 - acc: 0.9516  
Epoch 44/50  
124/124 [=====] - 0s - loss: 0.2249 - acc: 0.9516  
Epoch 45/50  
124/124 [=====] - 0s - loss: 0.2156 - acc: 0.9516  
Epoch 46/50  
124/124 [=====] - 0s - loss: 0.2070 - acc: 0.9516  
Epoch 47/50

```

124/124 [=====] - 0s - loss: 0.1987 - acc: 0.9516
Epoch 48/50
124/124 [=====] - 0s - loss: 0.1913 - acc: 0.9516
Epoch 49/50
124/124 [=====] - 0s - loss: 0.1838 - acc: 0.9597
Epoch 50/50
124/124 [=====] - 0s - loss: 0.1773 - acc: 0.9597

```

```
Out[31]: <tensorflow.python.keras._impl.keras.callbacks.History at 0x1aa6ae4d6d8>
```

```
In [32]: predictions = dnn_keras_model.predict_classes(scaled_x_test)
```

```
32/54 [=====>...] - ETA: 0s
```

```
In [33]: print(classification_report(predictions,y_test))
```

	precision	recall	f1-score	support
0	0.95	0.95	0.95	19
1	0.86	0.95	0.90	20
2	1.00	0.87	0.93	15
avg / total	0.93	0.93	0.93	54

## 5 Layers API

<https://www.tensorflow.org/tutorials/layers>

### 5.1 Formating Data

```

In [1]: import pandas as pd
        from sklearn.datasets import load_wine
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import MinMaxScaler

In [2]: wine_data = load_wine()
        feat_data = wine_data['data']
        labels = wine_data['target']

In [3]: X_train, X_test, y_train, y_test = train_test_split(feat_data,
                                                             labels,
                                                             test_size=0.3,
                                                             random_state=101)

```

```
In [4]: scaler = MinMaxScaler()
        scaled_x_train = scaler.fit_transform(X_train)
        scaled_x_test = scaler.transform(X_test)
        # ONE HOT ENCODED
        onehot_y_train = pd.get_dummies(y_train).as_matrix()
        one_hot_y_test = pd.get_dummies(y_test).as_matrix()
```

### 5.1.1 Parameters

```
In [6]: num_feat = 13
        num_hidden1 = 13
        num_hidden2 = 13
        num_outputs = 3
        learning_rate = 0.01

In [7]: import tensorflow as tf
        from tensorflow.contrib.layers import fully_connected
```

### 5.1.2 Placeholder

```
In [8]: X = tf.placeholder(tf.float32, shape=[None, num_feat])
        y_true = tf.placeholder(tf.float32, shape=[None, 3])
```

### 5.1.3 Activation Function

```
In [9]: actf = tf.nn.relu
```

### 5.1.4 Create Layers

```
In [10]: hidden1 = fully_connected(X, num_hidden1, activation_fn=actf)

In [11]: hidden2 = fully_connected(hidden1, num_hidden2, activation_fn=actf)

In [12]: output = fully_connected(hidden2, num_outputs)
```

### 5.1.5 Loss Function

```
In [13]: loss = tf.losses.softmax_cross_entropy(
        onehot_labels=y_true, logits=output)
```

### 5.1.6 Optimizer

```
In [14]: optimizer = tf.train.AdamOptimizer(learning_rate)
        train = optimizer.minimize(loss)
```



### 5.1.7 Init

```
In [15]: init = tf.global_variables_initializer()
```

```
In [17]: training_steps = 1000
        with tf.Session() as sess:
            sess.run(init)

            for i in range(training_steps):
                sess.run(train, feed_dict={X:scaled_x_train, y_true:onehot_y_train})

            # Get Predictions
            logits = output.eval(feed_dict={X:scaled_x_test})

            preds = tf.argmax(logits, axis=1)

            results = preds.eval()
```

```
In [18]: from sklearn.metrics import confusion_matrix, classification_report
        print(classification_report(results, y_test))
```

	precision	recall	f1-score	support
0	1.00	0.49	0.66	39
1	0.00	0.00	0.00	0
2	1.00	0.87	0.93	15
avg / total	1.00	0.59	0.73	54

```
C:\Users\Ripti\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1137: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to NaN because the truth labels contain zero positive samples. The precision value 0.00 was computed using the default setting of warn_for=('recall', 'true', 'average', warn_for)
```

## 6 TensorBoard

```
In [19]: import tensorflow as tf
        with tf.name_scope("OPERATION_A"):
            a = tf.add(1, 2, name="First_add")
            a1 = tf.add(100, 200, name='a_add')
            a2 = tf.multiply(a, a1)

        with tf.name_scope("OPERATION_B"):
            b = tf.add(3, 4, name='Second_add')
            b1 = tf.add(300, 400, name='b_add')
            b2 = tf.multiply(b, b1)
```

```

c = tf.multiply(a2,b2,name='final_result')

with tf.Session() as sess:
    writer = tf.summary.FileWriter("./output",sess.graph)
    print(sess.run(c))
    writer.close()

```

4410000

```

In [20]: k = tf.placeholder(tf.float32)

# Make a normal distribution, with a shifting mean
mean_moving_normal = tf.random_normal(shape=[1000], mean=(5*k), stddev=1)
# Record that distribution into a histogram summary
tf.summary.histogram("normal/moving_mean", mean_moving_normal)

# Setup a session and summary writer
with tf.Session() as sess:
    writer = tf.summary.FileWriter("./tmp/histogram_example")

    summaries = tf.summary.merge_all()

    # Setup a loop and write the summaries to disk
    N = 400
    for step in range(N):

        k_val = step/float(N)
        summ = sess.run(summaries, feed_dict={k: k_val})
        writer.add_summary(summ, global_step=step)

    writer.close()

```